


Lab 9 – Introduction to Program R and Distance Sampling for the Zombie Apocalypse.

Lab overview: This lab is designed to use the principles of distance sampling to evaluate the detection rate and density of zombies in a not-too-distant future post-apocalyptic scenario. The data you will be working with today were collected by survivors of the initial zombie virus outbreak, and consist of occurrence of zombies along two 1000 meter transect lines in a wooded landscape, not unlike that surrounding the University of Maine. Today we will analyze those data and derive estimates of zombie abundance and detection rates along the transect lines.

Recall that distance sampling is a technique for estimating animal (or zombie) abundance/density measured along transect lines or from fixed points. Observation of animals along the line (or at a point) are recorded along with a measure of their perpendicular distance from the line. We assume that animals on the line are detected perfectly ($p=1.0$), and that our ability to detect animals decreases as the animals are located further from the line. Based on this detection function we can correct for imperfect detection of individuals (right figure below), and estimate the proportion that were present but missed (Left figure – circles are detected animals, triangles are missed animals). Distance sampling is commonly used to estimate animal density (abundance/unit area) at different geographic locations or as a function of different habitat types.



Zombie Distance Analysis: On Brightspace you will find a file called Zdistance16.R. This file provides the completed code for analyzing the zombie distance sampling data. You can open the file in RStudio, so you will not need to type in code directly for the remainder of the lab. Notice that the file contains two pieces of information, first is the actual code to be run using R, and second is a series of comments I've made that

describe what each piece of code accomplishes. You will recognize the comments based on the double # that precedes them. As you know, this is a generic feature of R that allows you to include information in your code that is not actually read by the software. The steps I'll provide below will parallel the progression of the code, and in the code I'll add comments using a quadruple hashtag (####) when it's time to progress onto the next step.

1. First we need to input the distance data into R. You'll find the data saved as the csv file 'distall' on Brightspace. It includes two columns of information. The first delineates whether the transect line fell in deciduous or coniferous forest, and the second gives the distance for each observed zombie.
2. There are two options for importing the file to R. First you can use the code based on the `read.csv()` command provided in step 1 to import the data. To do so you will need to change the file path name to reflect the actual directory your file is stored in. Alternatively, use the "Import Dataset" option which can be accessed from the Tools menu or from the icon in the upper right window of R commander. If you go this route use the "From Text File" option, name the file "zDIST" (this is very important – if you name it something else the remainder of the code will need to be modified), and be sure to use the Headings=Yes option and to specify that the file uses a comma separator (note these should be defaults). Click Import.
3. Now we need to load the unmarked package. Do this by running the `install.packages()` and `library()` commands. You may be prompted after a second to select a "CRAN Mirror", which is just the server that the package will be pulled from. As far as I know it makes no difference which you select. I always choose the USA (CA1) mirror, which is housed at the University of California at Berkeley. If you are not asked to do this, just proceed with step 4.
4. Here, we're going to create a new field that defines a difference between the two transects and identifies that the new grouping variable "transect" as a factor. The following three lines of code will identify transect as a covariate that unmarked will recognize.
5. This chunk of code will create a data frame, `yDat`, which is formatted correctly for a distance sampling analysis. When you look at it you'll notice the observations have been pooled into 5 discrete bins in 10 m increments, and they are separated among the two transects.

6. Using the `yDat` and `covs` data frames, we'll construct an input file that will contain all the data for our distance sampling analysis and will contain some other relevant information like transect length and the fact that this was a line-based sample, as opposed to point-based. I've also provided some code that will produce a histogram of your data to show how your detections varied based on distance from the transect line. Recall that in a distance analysis we assume that animal (zombie) distribution is random with respect to the line, but that our ability to detect the zombies decreases as we move away from the line. Hopefully your histogram is consistent with the later assumption.

Also notice that when you create these figures in the plot window you can export them as image files using the "Export" feature.

7. Our next step will be to identify the best distance-detection function, given your data. Again, we assume that detection declines with distance from the line, however, the specific way that detection changes with distance can be variable among studies for a variety of reasons. For example, detection may be high very close to the line but then may very rapidly decline. This might occur in very dense vegetation or in highly variable terrain. In contrast detection may decline slowly and progressively with distance, for example in even terrain or with more moderate vegetation. There are three general distance-detection functions we can use to approximate this: the half normal, hazard, and negative exponential functions. If you run the code associated with step #7, you can see what each of these three functions look like.

Notice that these three functions require different parameters. The half-normal is, like the name implies, half of a normal distribution which can be characterized by a single parameter, sigma (or the variance of the distribution). The hazard function requires two parameters that define the shape and scale of the curve, and the negative exponential requires one parameter describing the rate of decline. For the purpose of these examples I've defined these all arbitrarily with a value of 10 for each respective parameter.

8. To test which of the three distance-detection functions are most appropriate given our data, we can run three null models and only change the shape of the distance-detection function. This will also be your first chance to run a distance model. Take a quick look at the code. Importantly, the `~1 ~1` component is where we define the structure of the model, and `"1"` defines an intercept-only model. The first term is associated with the detection parameter, and the second associated with the density parameter. Here we are creating three objects, each of them distance sampling models with an intercept-only structure on both

detection and density parameters, where the only difference is the shape of the distance-detection function. *Note – if you get an error message saying that the Hessian is singular, you can ignore it and move on.

9. Now, using the same principles of AIC model selection that we've used in past weeks, we can compare the relative fit of our three alternative models, given our data. The first batch of code here combines the model results into a single object, `null.fit`, and the next three statements produce an AIC table, provides us with the parameter estimates from each model, and the third gives the standard errors for the estimates. You will want to check two different things with these results. First, use AIC to determine which model provides the best fit to the data. Second check the standard errors for each model to see that the parameters are actually being estimated. A returned value of "NaN" indicates convergence was not reached on the parameter, suggesting the data were insufficient to fit the model in general. Notice that "NA" means something different – that the specific parameter is not relevant to the given model. The "best" model will be the one that produces the lowest AIC score while also estimating all necessary parameters. Based on these results, determine which function best fits your specific data.
10. After finding the best distance-detection function, the next step will be to test some other alternative models. Really here we only have one other variable to consider, and that is the difference among the two transects that you ran (which also reflects the difference among the two observers). The four lines of code given here each include a different combination of either running intercept-only structure to detection or density, or allowing detection and/or density to differ among the two transects. Remember that the model structure is given by the notation `~ detection ~ density`, where `~1` reflects a null model and `~Trans` gives a model where either detection or density is allowed to differ among the two transects. To come full circle to previous weeks, the linear model for a `~Trans` structure to the detection function would look something like the following:

$$Detection = \beta_0 + \beta_1 * Transect$$

where Beta 0 would be the model intercept, and Beta 1 is the effect of transect number on detection. Thus, when calling the `summary()` function on the model `p.Trans`, the `sigma(Intercept)` Estimate is the value for Beta 0 (model intercept) and the `sigmaTrans2` Estimate value provides the effect of being Transect 2, where Transect 1 is defined strictly by the intercept (remembering that Transect Number is a factor, making it equivalent to a dummy variable in our prior weeks).

Based on your assessment of the AIC results and coefficient estimates from this 3-model set, determine whether there were differences in either zombie detection or density between the two transect lines, and what those differences were (i.e. which transect was higher or lower in detection or density, if at all).

11. A major objective of a distance analysis is to actually generate an estimate of density (animals per unit area) and abundance (total number along the transect line). For our distance analysis in unmarked, our actual response variable is density, and that is what our model results report. To generate a density estimate, we will use the `backTransform()` function, which will produce and report the density estimate. In this case if we use the Null model density is averaged across the two transect lines.
12. Next we can convert density into an actual estimate of the total number of zombies found along the transect lines. This will require doing a small amount of math to take the estimate of density and convert it to abundance given the size of the area we sampled. Remember the data are based on two 1000 meter transects, and our maximum distance was 50 m from the line. However, to compute the actual abundance of zombies we've got to first compute the effective transect width. This uses the distance-detection function to compute the actual functional area that we covered during the survey, accounting for the fact that detection declines as we move further away from the survey line. The 'integrate' function is what computes the area under the detection curve by, you guessed it, integrating the half-normal function.
13. Based on the effective width we can convert density to an estimate of abundance, and the second batch of code here will use bootstrapping to derive a standard error for that estimate.
14. The last batch of code will cover how to plot the detection estimates to produce a distance-detection graph such as that contained on the first page of the handout.

Assignment. We've just ran a fairly involved Distance Sampling analysis in program R. From it, you should have assessed the appropriate distance-detection function, evaluated support for differences among the two transect lines, estimated the abundance of zombies along the transect lines, and plotted your detection functions. We also covered the basic fundamentals of working in program R. Based on what you've learned during the lab, answer/complete the following questions and turn them in as the assignment for this week.

1. Which distance-detection function did you find was best supported by the data? What does the shape of this function imply about your ability to detect zombies in the Maine woods? Please include with this answer the AIC table* that summarized the results of your three distance-detection models, and your final plot** of detection by distance with the overlaid distance-detection function (i.e. the graph from step 14).
2. A colleague of yours working in a major urban area collected similar data, and found that a hazard function with a shape parameter=3.9 and a scale parameter =8.2 best fit the zombie distance-detection relationship in her system. What does this imply about zombie detection in an urban setting versus what we found in a forest environment? How would that information be relevant for survival during a zombie viral outbreak?
3. Did you find any differences among the two transect lines with respect to density or detection? If so, what were those differences? Use supporting information from your AIC table or parameter estimates to support your answer, and also include your final AIC table (the second one from the analysis).
4. What was your final estimate of zombie abundance? What was your estimate of zombie density? If we surveyed a 100 ha woodlot, what would you predict was the total abundance of zombies within the entire woodlot? What inherent assumption are you making when extrapolating the estimates from your two transect lines to the woodlot as a whole?

* Notice that AIC tables can be copied out of R into excel using the `write.csv()` command. If you type `?write.csv` as an R command line, you will find instructions for exporting using this command, or look back to previous labs. The same general formatting rules we've used all semester apply.

** Graphs can be saved as .jpeg files using the "Export" feature in the graphing window.

